

JOEL IVÁN SOLÍS

40 Boston Ave ◊ Somerville, MA 02144 ◊ (956) · 459 · 4809

✉ jsolis@alum.mit.edu ◊    jsolis1199

SKILLS

CI/CD	GitHub Actions, GitLab CI	Warehouses	Redshift, Snowflake, BigQuery
DevOps	Docker, Terraform	Orchestration	Airflow, MWAA
Formats	Parquet, Iceberg	Ingestion	Fivetran
Languages	Python, PL/SQL, Java	Processing	dbt, Glue, EMR

EXPERIENCE

Tinuiti Apr 2022 - Present
Data Engineer II *Remote*

- Designed a CI/CD framework using DynamoDB, GitHub Actions, and a Fivetran API wrapper written in Python to streamline creation and management of our S3 data lake connectors
- Migrated ingestion of Facebook, LinkedIn, Search Ads 360, and Snapchat ad data from a legacy ETL provider to Fivetran, which allowed our CTO to cancel the previous contract and ultimately save Tinuiti millions of dollars per annum
- Improved the reliability of an error-prone data pipeline handling encrypted first-party data by refactoring its decryption logic and porting it to Amazon Glue, resulting in Redshift loads finishing 75% faster and job failures occurring less than once a year
- Modernized Tinuiti's underlying data platform by moving analytic logic out of costly and unmaintainable Tableau Prep Flows and into a business-conformed data mart in dbt

Datto Oct 2020 - Mar 2022
Data Engineer II *Remote*

- Migrated our data warehouse from an on-premises monolithic MySQL instance, Spark cluster, and Airflow scheduler to a cloud implementation on AWS using Snowflake, EMR, and MWAA
- Bolstered the engineering capabilities of the Data & Analytics team by designing a dbt development workflow that enables analysts without software engineering backgrounds to version their SQL, test the code, and submit their work for production deployment all within VS Code

Insight Data Science Jan 2020 - Jul 2020
Program Director, Data Engineering *Boston, MA*

- Introduced a cohort of 17 former academics into the data engineering industry by teaching them about distributed computing frameworks and provided mentorship as they developed their own scalable data pipeline in 3 short weeks
- Deployed an 11-node arbitrage monitoring cluster on AWS using Kafka, Spark, and Cassandra to stream live cryptocurrency trades from the WebSocket endpoints of 5 different cryptocurrency exchanges

EDUCATION

Massachusetts Institute of Technology Jun 2014 - Jun 2018
B.S. in Mathematics & Physics